

# *Secondary Storage*

- ❑ External Sort

- ❑ B-tree Index

- ❑ Other Indices

# Secondary Storage: *Basics*

## □ Sequential Access vs. Direct Access (*Random Access*)

- Block access → organize files as user-defined blocks

## □ File manger in OS supports ...

- *Cluster*: a number of contiguous sectors
- Once a cluster has been found on a disk, all sectors in that cluster can be accessed without an additional seek.
- If a file consists entirely of contiguous clusters, the seeking time is minimized.
- As the number of “*non-contiguous*” clusters in a file increases, the file becomes more spread out on the disk, and the amount of seeking necessary increases.

# Secondary Storage: *A Journey of A Byte*

*What happens when `write(textfile, 'P', 1)` is executed?*

Part that takes place *in memory*:

- ❑ Call **OS** to oversee this operation

- ❑ Call **File manager** (*A part of OS*)

- Check whether the operation is permitted
  - **the file is open, the type of access is allowed, ...**
- Locate the **physical location** where the byte will be stored (*Drive, Cylinder, Track & Sector*)
- Make sure that the **sector** to locate the 'P' is already in a **system I/O Buffer** and deposit 'P' into it
- Call I/O processor to send the sector back to disk

# Secondary Storage: *A Journey of A Byte*

Part that takes place *outside* of memory:

## □ I/O Processor

- Wait for an **external data path** to become available (CPU is faster...)

### ■ *Direct Memory Access* Input/Output

## □ Disk Controller

- I/O Processor asks the disk controller whether the disk drive is available for writing
- Disk Controller instructs the disk drive to move its read/write head to the right **track** (**seek time**) and then wait for the desired **sector** (latency time)
- Disk spins to right location and **'P'** is written (transfer time)

# Secondary Storage: *External Sort*

- Internal Sort: main memory
- External Sort: secondary storage + main memory
- Example

- How to sort 6,400 student records? (if each record takes 100KB, we need 640MB in total)
- Can we use **very less space** to do the same job?

- The records are divided into 64 blocks


- Each block keeps only 100 student records (10MB)

Step 1. *Internal sort on each block*

Step 2. (**external**) *Merge sort*

# Secondary Storage: *External Sort*

## Step 1. Internal sort



Memory space: 10MB

$R_{64,1}, R_{64,2}, \dots, R_{64,100}$


$R_{1,1}, R_{1,2}, \dots, R_{1,100}$

$R_{2,1}, R_{2,2}, \dots, R_{2,100}$

...


$R_{64,1}, R_{64,2}, \dots, R_{64,100}$

## Step 2. Merge sort



Memory space: 40MB

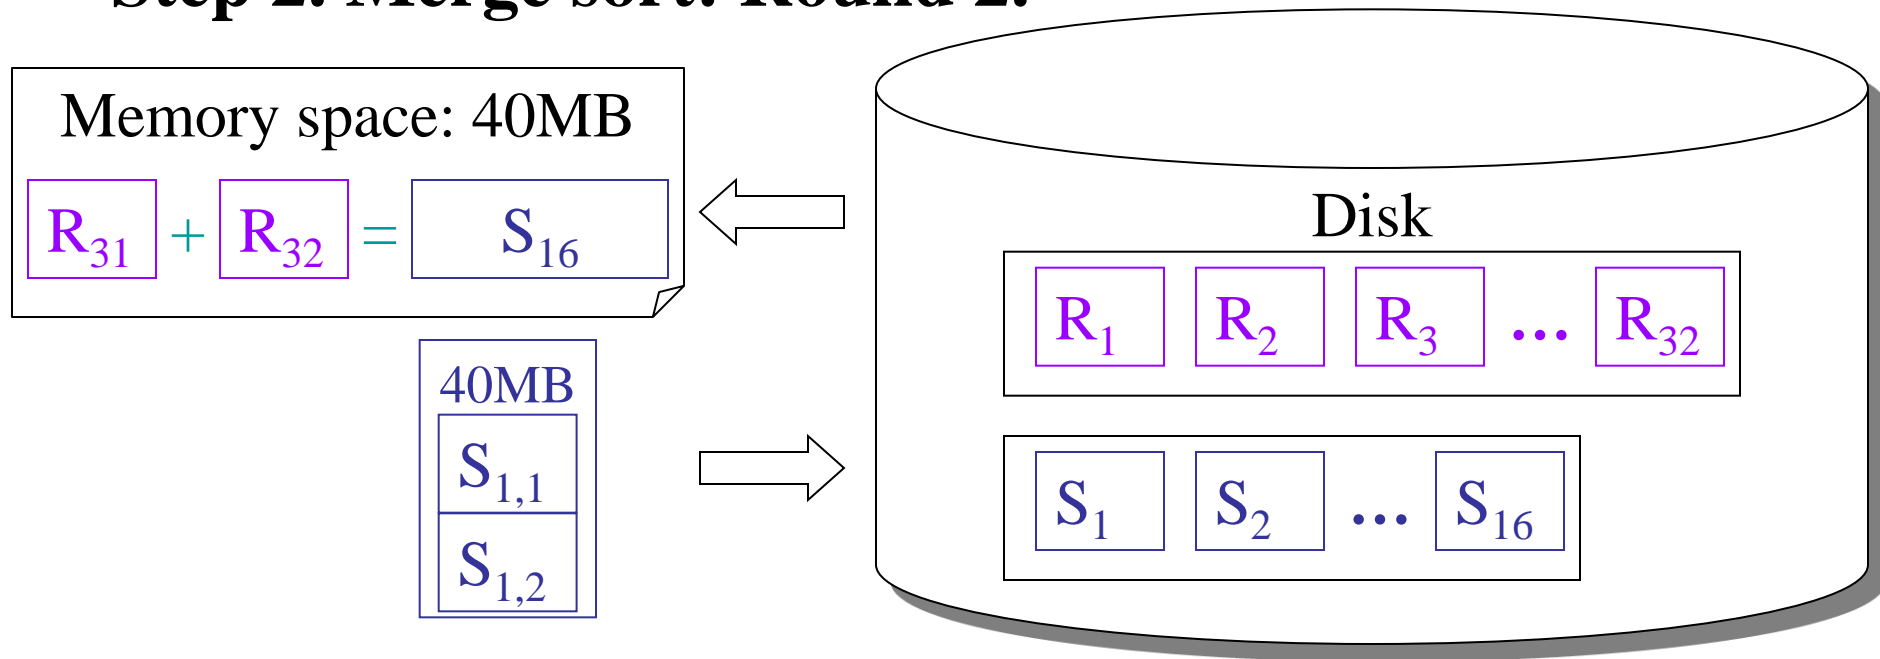
$B_{63} + B_{64} = R_{32}$




$R_1, R_2, R_3, \dots, R_{32}$

# Secondary Storage: *External Sort*

## Step 2. Merge sort: Round 2.



## Step 2. Merge sort: Round 3.





# Secondary Storage: *External Sort*


## Step 2. Merge sort: Round 4.

Memory space: 40MB

$$B_1 + B_2 = R_1$$



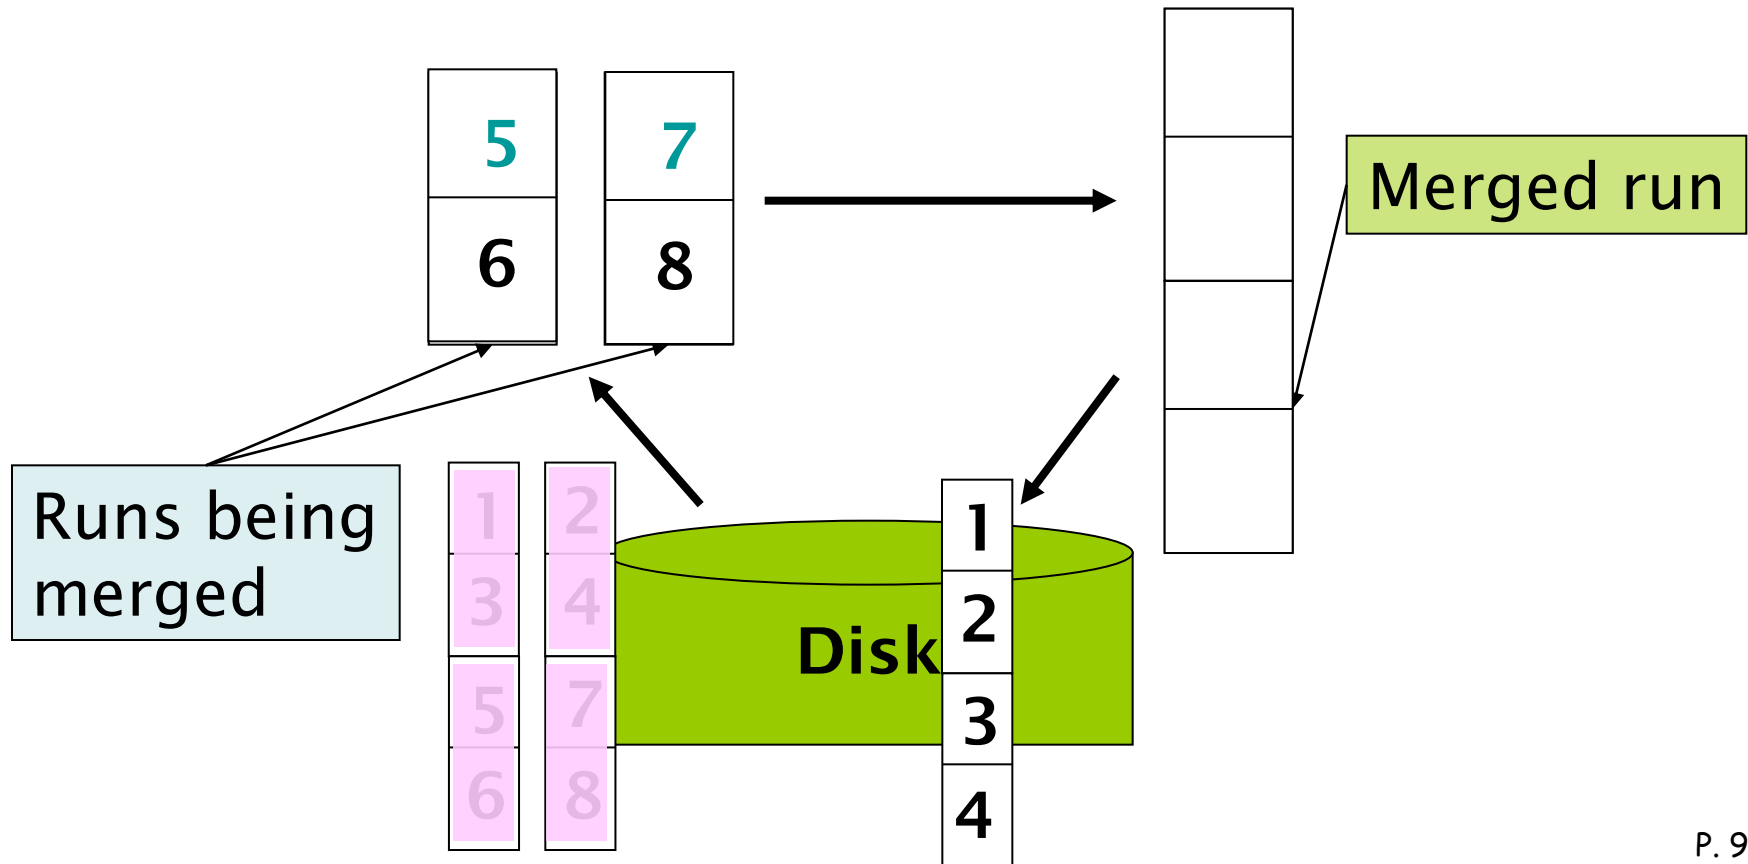
## Step 2. Merge sort: Rounds 5 & 6.





# External Sort: 2-way Merge

- ❑ Input buffer
- ❑ Output buffer



# External Sort: *K-way Merge*

## □ 2-way merge

- 64 runs  $\rightarrow 32, 16, 8, 4, 2, 1 \rightarrow \log_2 64 = 6$  passes
- 16 runs  $\rightarrow 8, 4, 2, 1 \rightarrow \log_2 16 = 4$  passes

□ A **k-way** merge on **m runs** needs  $\log_k m$  passes

□ **Higher-order** merge can reduce I/O time


## □ 4-way merge

- 64 runs  $\rightarrow 16, 4, 1 \rightarrow \log_4 64 = 3$  passes
- 16 runs  $\rightarrow 4, 1 \rightarrow \log_4 16 = 2$  passes

6	7	9	8
⋮	⋮	⋮	⋮


# K-way Merge: *Selection Tree*

- Given  $k$  sorted runs to be merged, a data structure named *selection tree* can reduce the number of comparisons for *finding the next smallest element*.



# K-way Merge: *Selection Tree*

## □ Selection tree on $k$ sorted runs



# Secondary Storage: *File Structures*

## □ Record

### – Field

■ name

■ value

Last Name: 'Jordan'  
First Name: 'Michael'  
Score: 30  
Assist: 5  
Rebound: 6  
Champion: 6  
MVP: 5

Last Name: 'Chamberlain'  
First Name: 'Wilt '  
Score: 30  
Assist: 4  
Rebound: 22  
Champion: 2  
MVP: 4

Last Name: 'Johnson'  
First Name: 'Earvin'  
Score: 20  
Assist: 11  
Rebound: 7  
Champion: 5  
MVP: 3

Last Name: 'Abdul-Jabbar'  
First Name: 'Kareem '  
Score: 25  
Assist: 4  
Rebound: 11  
Champion: 6  
MVP: 6

# File Structures: *Fields*

□ Many ways of adding structure to files to maintain the identity of *fields*:

1. Force the field into a predictable length
2. Begin each field with a length indicator
3. Separate the fields with delimiters
4. Use a “*fieldname = value*” expression to identify each field and its content.

■ *Self-describing*

Jordan	← unused space →
Michael	← unused space →
...	

Q&A: Which is the best choice if some fields are missed?

# File Structures: *Records*

## □ Methods to organize the records of a file:

1. Requiring that the records be a **predictable number of *bytes*** in length.
2. Requiring that the records be a **predictable number of *fields*** in length.
3. Beginning each record with a **length indicator** consisting of a count of the number of bytes that the record contains.

Jordan,Michael,30,5,6,6,5

← unused space →

Johnson,Earvin,20,11,7,5,3

← unused space →

...

# File Structures: *Records*

## □ Methods to organize the records of a file:

4. Placing a **delimiter** at the end of each record to separate it from the next record.

Jordan|Michael|30|5|6|6|5|#Johnson|Earvin|20|11|7|5|3|#...

5. Using a **second file (index)** to keep track of *the beginning byte address for each record*.

4 Jordan|Michael|30|5|6|6|5|Johnson|Earvin|20|11|7|5|3|...





# Variable-length vs. Fixed-length

□ Offset of fixed-length record  $RRN = i$

(Header record length) +  $(i - 1) * (\text{record length})$


$RRN=2$ : offset =  $1 + (2-1) * 100 = 101$  bytes



**RRN: relative record number**

[1]	Jordan,Michael,30,5,6,6,5	← unused space →
[2]	Johnson,Earvin,20,11,7,5,3	← unused space →
...	...	

4	Jordan Michael 30 5 6 6 5 Johnson Earvin 20 11 7 5 3 ...
---	--



# Data Structures

P. 18


# Fixed-length Records: *Free List*

## □ List head

- Stored in the file header
- Keep the *RRN* of one *deleted record*

## □ Use one field of the deleted record to keep the *RRN* of the next deleted record

## □ Regard these *RRNs* (offset) as **pointers** in the file



# Key Sort

- ❑ When sorting a file in memory, the only thing that really needs sorting are the **keys** of records.
- ❑ Key sort algorithms work like *internal sort*, but with 2 important differences:
  - Rather than read an **entire record** into a memory array, we simply read each record into a **temporary buffer**, *extract the key and then discard the record*.
  - If we want to write the records in sorted order, we have to *read them one more time*.

# Key Sort: *Pseudocode*

```
Op  
Cre  
Re  
REC_COUNT = record count stored in head record  
// read all the records in sequence  
for i = 1 to REC_COUNT  
    do loop 1  
    sort (KEY_ARRAY)  
// repeatedly  
    for i = 1 to REC_COUNT  
        do loop 2  
Close IN_FILE and OUT_FILE
```

1. read record from IN\_FILE into BUFFER (*in order*)
2. KEY\_ARRAY [i].KEY = extract **key** from BUFFER
3. KEY\_ARRAY [i].RRN = i

1. j = KEY\_ARRAY [i].RRN
2. **seek** in IN\_FILE to the record with RRN=j
3. read record from IN\_FILE into BUFFER
4. write BUFFER to OUT\_FILE (*in order*)

# Illustration I: *Key Sort*

(KEY RRN)

(99, 1)

(95, 2)

(88, 3)

(99, 4)

(95, 5)

(KEY RRN)

(99, 1)

(99, 4)

...

(95, 2)

(95, 5)

...

(88, 3)

...

REC\_COUNT (number of records)

record 1: BS001, 99, Lee, ...

record 2: BS003, 95, Lin, ...

record 3: BS004, 88, Wang, ...

record 4: BS005, 99, Liu, ...

record 5: BS008, 95, Chen, ...

record 1: BS001, 99, Lee, ...

record 4: BS005, 99, Liu, ...

...

record 2: BS003, 95, Lin, ...


IN\_FILE

OUT\_FILE

# Key Sort: *Limitation*

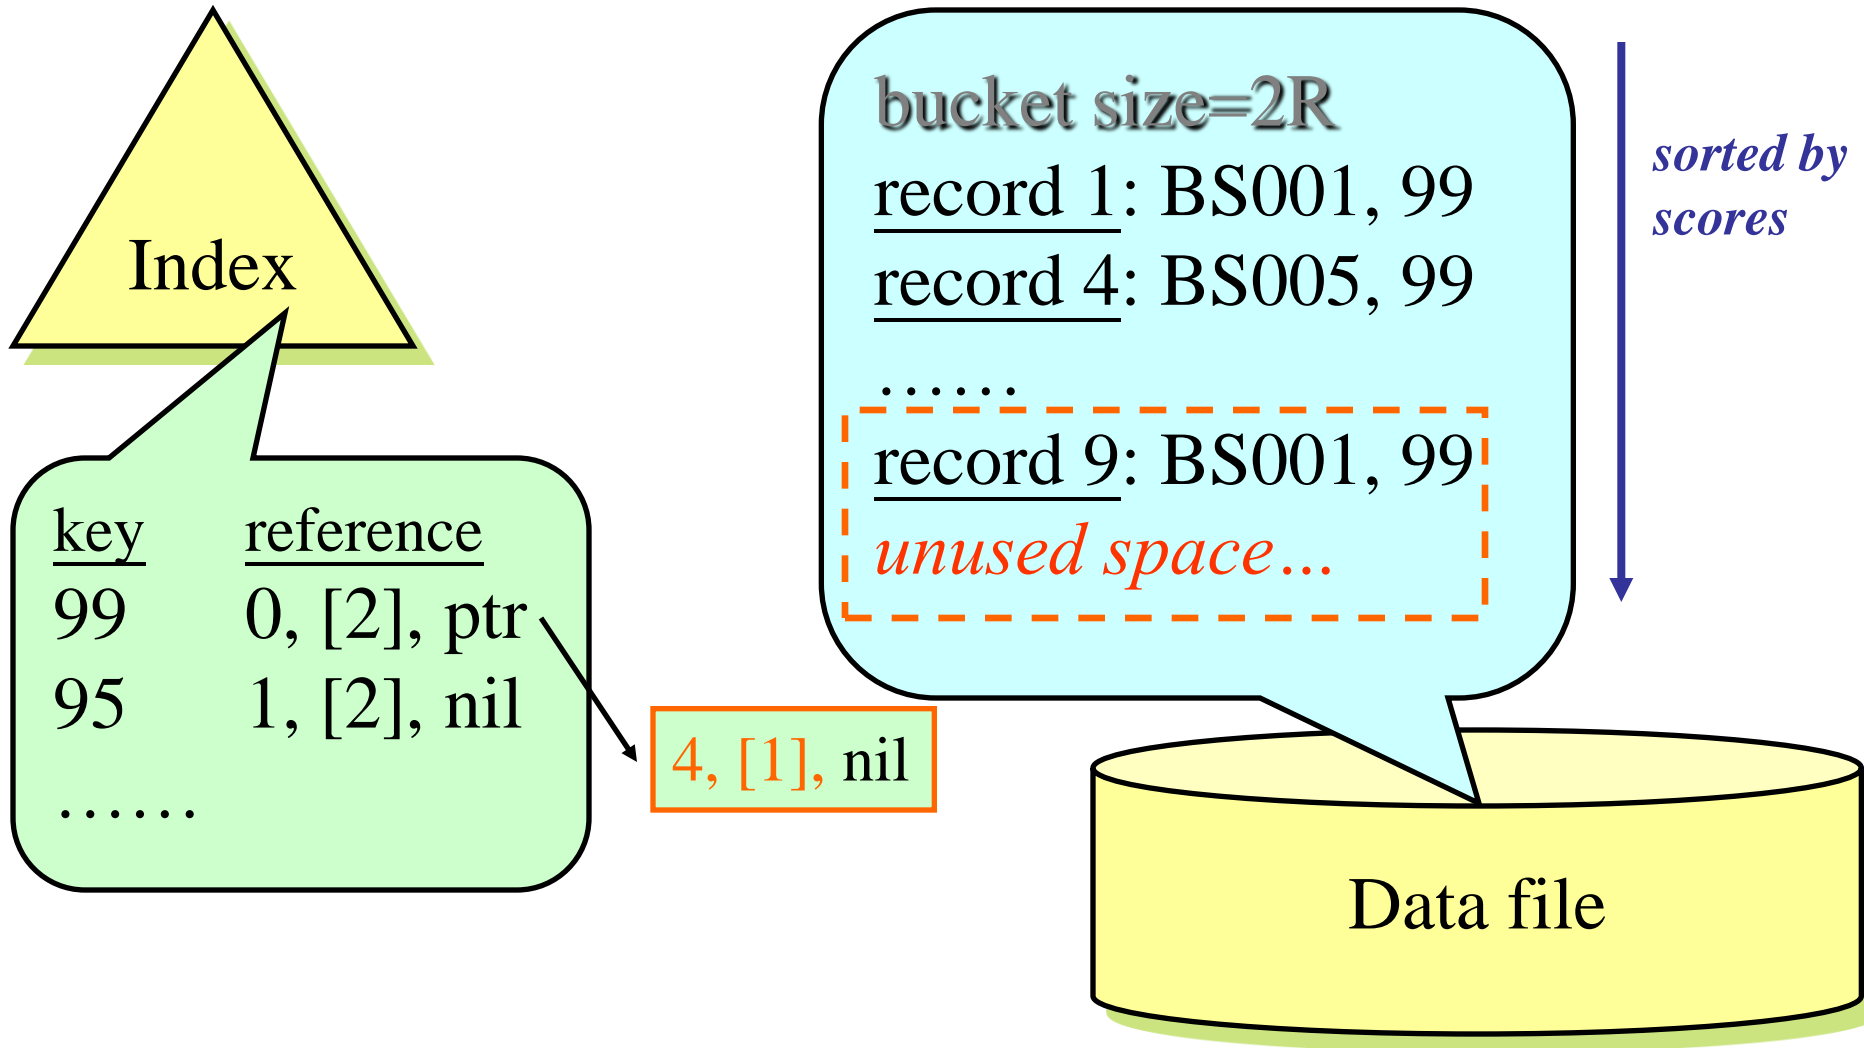
- ❑ Writing the records in sorted order requires as many **random seeks** as there are records.
- ❑ Since writing is interspersed with reading, writing also requires as many seeks as there are records.
- ❑ Why bother to write **the file of records** in the order of their keys?
  - *Simply write back the sorted index!*

# Illustration II: *Secondary Index*





# Illustration IV: *Bucket Chaining*



# Illustration V: *Backup & Reload*

Index

<u>key</u>	<u>reference</u>
99	0, [2], ptr
95	1, [2], nil
.....	

4, [1], nil

Data file

99, 0, 2  
99, 4, 1  
95, 1, 2  
.....

Index (*backup*) file